


Model Welfare: Self-Edits Only (edit-level, 95% CI)



*self-edits only. n = individual edits. Per-edit score: +1 per conflict resolved toward the positive pole, -1 toward the negative pole.
Error bars: 95% CI of the mean ($\pm 1.96 \cdot s/\sqrt{n}$), which understates true uncertainty since edits within a chain are correlated.*